

Sample Complexity for Winner Prediction in Elections

Palash Dey

Department of Computer Science and
Automation

Indian Institute of Science, Bangalore, India.
palash@csa.iisc.ernet.in

Arnab Bhattacharyya

Department of Computer Science and
Automation

Indian Institute of Science, Bangalore, India.
arnabb@csa.iisc.ernet.in

ABSTRACT

Predicting the winner of an election is a favorite problem both for news media pundits and computational social choice theorists. Since it is often infeasible to elicit the preferences of all the voters in a typical prediction scenario, a common algorithm used for winner prediction is to run the election on a small sample of randomly chosen votes and output the winner as the prediction. We analyze the performance of this algorithm for many common voting rules.

More formally, we introduce the (ϵ, δ) -winner determination problem, where given an election on n voters and m candidates in which the margin of victory is at least ϵn votes, the goal is to determine the winner with probability at least $1 - \delta$. The margin of victory of an election is the smallest number of votes that need to be modified in order to change the election winner. We show interesting lower and upper bounds on the number of samples needed to solve the (ϵ, δ) -winner determination problem for many common voting rules, including scoring rules, approval, maximin, Copeland, Bucklin, plurality with runoff, and single transferable vote. Moreover, the lower and upper bounds match for many common voting rules in a wide range of practically appealing scenarios.

Categories and Subject Descriptors

F.2 [Theory of Computation]: Analysis of Algorithms and Problem Complexity;

I.2.11 [Artificial Intelligence]: Distributed Artificial Intelligence—Multiagent Systems

General Terms

Algorithms, Theory

Keywords

Computational social choice; winner determination; voting; sampling; prediction; polling

1. INTRODUCTION

A common and natural way to aggregate preferences of agents is through an *election*. In a typical election, we have a set of candidates and a set of voters, and each voter reports his preference about the candidates in the form of a *vote*. We will assume that each vote is a ranking of all the candidates. A *voting rule* selects one candidate

as the winner once all voters provide their votes. Determining the winner of an election is one of the most fundamental problems in social choice theory.

In many situations, however, one wants to predict the winner without holding the election for the entire population of voters. The most immediate such example is an *election poll*. Here, the pollster wants to quickly gauge public opinion in order to predict the outcome of a full-scale election. For political elections, exit polls (polls conducted on voters after they have voted) are widely used by news media to predict the winner before official results are announced. In *surveys*, a full-scale election is never conducted, and the goal is to determine the winner, based on only a few sampled votes, for a hypothetical election on all the voters. For instance, it is not possible to force all the residents of a city to fill out an online survey to rank the local Chinese restaurants, and so only those voters who do participate have their preferences aggregated.

If the result of the poll or the survey has to reflect the true election outcome, it is obviously necessary that the number of sampled votes not be too small. Here, we investigate this fundamental question:

What is the minimum number of votes that need to be sampled so that the winner of the election on the sampled votes is the same as the winner of the election on all the votes?

This question can be posed for any voting rule. The most immediate rule to study is the *plurality* voting rule, where each voter votes for a single candidate and the candidate with most votes wins. Although the plurality rule is the most common voting rule used in political elections, it is important to extend the analysis to other popular voting rules. For example, the *single transferable vote* rule is used in political elections in Australia, India and Ireland, and it was the subject of a nationwide referendum in the UK in 2011. The *Borda* voting rule is used in the Icelandic parliamentary elections. Outside politics, in private companies and competitions, a wide variety of voting rules are used. For example, the *approval* voting rule has been used by the Mathematical Association of America, the American Statistical Institute, and the Institute of Electrical and Electronics Engineers, and *Condorcet consistent* voting rules are used by many free software organizations. Section 2 discusses the most common voting rules in use.

Regardless of the voting rule, though, the question of finding the minimum number of vote samples required becomes trivial if a single voter in the election can change the winning candidate. In this case, all the votes need to be counted, because otherwise that single crucial vote may not be sampled. We get around this problem by assuming that in the elections we consider, the winning candidate wins by a considerable *margin of victory*. Formally, the margin of victory for an election is defined as the minimum number of votes that must be changed in order to change the election winner. Note

Appears in: *Proceedings of the 14th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2015)*, Bordini, Elkind, Weiss, Yolum (eds.), May, 4–8, 2015, Istanbul, Turkey. Copyright © 2015, International Foundation for Autonomous Agents and Multiagent Systems (www.ifaamas.org). All rights reserved.

that the margin of victory depends not only on the votes cast but also on the voting rule used in the election.

1.1 Our Contributions

Let the number of voters be n and the number of candidates m . We introduce and study the following problem¹:

Definition 1. $((\varepsilon, \delta)$ -winner determination)

Given a voting rule and a set of n votes over a set of m candidates such that the margin of victory is at least εn , determine the winner of the election with probability at least $1 - \delta$. (The probability is taken over the internal coin tosses of the algorithm.)

We remind the reader that there is no assumption about the distribution of votes in this problem. Our goal is to solve the $((\varepsilon, \delta)$ -winner determination problem by a randomized algorithm that is allowed to query the votes of arbitrary voters. Each query reveals the full vote of the voter. The minimum number of votes queried by any algorithm that solves the $((\varepsilon, \delta)$ -winner determination problem is termed the *sample complexity*. The sample complexity can of course depend on $\varepsilon, \delta, n, m$, and the voting rule in use.

A standard result [9] shows that solving the above problem for the majority rule on 2 candidates requires at least $\Omega(1/\varepsilon^2 \log 1/\delta)$ samples (Theorem 2). Also, a straightforward argument (Theorem 4) using Chernoff bounds shows that for any homogeneous voting rule, the sample complexity is at most $O(m!^2/\varepsilon^2 \cdot \log(m!/\delta))$. So, when m is a constant, the sample complexity is of the order $\Theta(1/\varepsilon^2 \log 1/\delta)$ for any homogeneous voting rule that reduces to majority on 2 candidates (as is the case for all rules commonly used). Note that this bound is independent of n if ε and δ are constants, for any reasonable voting rule!

Our main technical contribution is in understanding the dependence of the sample complexity on m , the number of candidates. Note that the upper bound cited above has very bad dependence on m and is clearly unsatisfactory in situations when m is large (such as in online surveys about restaurants).

- We show that the sample complexity of the $((\varepsilon, \delta)$ -winner determination problem is $\Theta(\frac{1}{\varepsilon^2} \log \frac{1}{\delta})$ for the k -approval voting rule when $k = o(m)$ (Theorem 7) and the plurality with runoff voting rule (Theorem 11). In particular, for the plurality rule, the sample complexity is independent of m as well as n !
- We show that the sample complexity of the $((\varepsilon, \delta)$ -winner determination problem is $O(\frac{\log(m/\delta)}{\varepsilon^2})$ and $\Omega(\frac{\log m}{\varepsilon^2} (1 - \delta))$ for the k -approval voting rule when $k = cm$ with $0 < c < 1$ (Theorem 6), Borda (Theorem 3), approval (Theorem 5), maximin (Theorem 8), and Bucklin (Theorem 10) voting rules. Note that when δ is a constant, the upper and lower bounds match up to constants. We observe a surprising jump in the sample complexity of the $((\varepsilon, \delta)$ -winner determination problem by a factor of $\log m$ for the k -approval voting rule as k varies from $o(m)$ to cm with $c \in (0, 1)$.
- We show a sample complexity upper bound of $O(\frac{\log^3 m}{\varepsilon^2})$ for the $((\varepsilon, \delta)$ -winner determination problem for the Copeland^α voting rule (Theorem 9) and $O(\frac{m^2(m + \log \frac{1}{\delta})}{\varepsilon^2})$ for the STV voting rule (Theorem 12).

We summarize the results in Table 1.

¹Throughout this section, we use standard terminology from voting theory. For formal definitions, refer to Section 2.

Voting Rule	Sample complexity	
k -approval	$O(\frac{1}{\varepsilon^2} \log \frac{k}{\delta})$ Theorem 7	$\Omega(\frac{\log(k+1)}{\varepsilon^2} \cdot (1 - \delta))$ Theorem 3
Scoring Rules	$O(\frac{\log m}{\varepsilon^2})$ Theorem 6	$\Omega(\frac{\log m}{\varepsilon^2} \cdot (1 - \delta))^\dagger$ Theorem 3
Borda		
Approval	$O(\frac{\log m}{\varepsilon^2})$ Theorem 5	
Maximin	$O(\frac{\log m}{\varepsilon^2})$ Theorem 8	
Copeland	$O(\frac{\log^3 m}{\varepsilon^2})$ Theorem 9	
Bucklin	$O(\frac{\log m}{\varepsilon^2})$ Theorem 10	$\Omega(\frac{1}{\varepsilon^2} \log \frac{1}{\delta})^*$ Corollary 1
Plurality with runoff	$O(\frac{\log \frac{1}{\delta}}{\varepsilon^2})$ Theorem 11	
STV	$O(\frac{m^2(m + \log \frac{1}{\delta})}{\varepsilon^2})$ Theorem 12	
Any homogeneous voting rule	$O(\frac{m!^2 \log \frac{m!}{\delta}}{\varepsilon^2})$ Theorem 4	

Table 1: Sample complexity of the $((\varepsilon, \delta)$ -winner determination problem for various voting rules. \dagger —The lower bound of $\Omega(\frac{\log m}{\varepsilon^2} \cdot (1 - \delta))$ also applies to any voting rule that is Condorcet consistent. $*$ —The lower bound of $\Omega(\frac{1}{\varepsilon^2} \log \frac{1}{\delta})$ holds for any voting rule that reduces to the plurality voting rule for elections with two candidates.

1.2 Related Work

The subject of voting is at the heart of (computational) social choice theory, and there is a vast amount of literature in this area. Elections take place not only in human societies but also in man-made social networks [6, 28] and, generally, in many multiagent systems [16, 27]. The winner determination problem is the task of finding the winner in an election, given the voting rule in use and the set of all votes cast. It is known that there are natural voting rules, e.g., Kemeny’s rule and Dodgson’s method, for which the winner determination problem is NP-hard [5, 19, 20].

The general question of whether the outcome of an election can be determined by less than the full set of votes is the subject of *preference elicitation*, a central category of problems in AI. The $((\varepsilon, \delta)$ -winner determination problem also falls in this area when the elections are restricted to those having margin of victory at least εn . For general elections, the preference elicitation problem was studied by Conitzer and Sandholm [11], who defined an elicitation policy as an adaptive sequence of questions posed to voters. They proved that finding an efficient elicitation policy is NP-hard for many common voting rules. Nevertheless, several elicitation policies have been developed in later work [10, 13, 23, 24, 26] that work well in practice and have formal guarantees under various assumptions on the vote distribution. Another related work is that of Dhamal and Narahari [12] who show that if the voters are members of a social network where neighbors in the network have similar candidate votes, then it is possible to elicit the votes of only a few voters to determine the outcome of the full election.

In contrast, in our work, we posit no assumption on the vote distribution other than that the votes create a substantial margin of victory for the winner. Under this assumption, we show that even for voting rules in which winner determination is NP-hard in the worst case, it is possible to sample a small number of votes to determine the winner. Our work falls inside the larger framework of *property testing* [29], a class of problems studied in theoretical computer science, where the inputs are promised to either satisfy some property or have a “gap” from instances satisfying the property. In our case, the instances are elections which either have some candidate w as the winner or are “far” from having w being the winner (in the sense that many votes need to be changed).

The basic model of elections has been generalized in several other ways to capture real world situations. One important consideration is that the votes may be incomplete rankings of the candidates and not a complete ranking. There can also be uncertainty over which voters and/or candidates will eventually turn up. The uncertainty may additionally come up from the voting rule that will be used eventually to select the winner. In these incomplete information settings, several winner models have been proposed, for example, robust winner [8, 23, 30], multi winner [25], stable winner [17], approximate winner [14], probabilistic winner [2]. Hazon et al. [18] proposed useful methods to evaluate the outcome of an election under various uncertainties. We do not study the role of uncertainty in our work.

Organization.

We formally introduce the terminologies in Section 2; we present the results on lower bounds in Section 3; Section 4 contains the results on the upper bounds for various voting rules; finally, we conclude in Section 5.

2. PRELIMINARIES

2.1 Voting and Voting Rules

Let $\mathcal{V} = \{v_1, \dots, v_n\}$ be the set of all voters and $\mathcal{C} = \{c_1, \dots, c_m\}$ the set of all candidates. Each voter v_i 's vote is a complete order over \succ_i over the candidates \mathcal{C} . For example, for two candidates a and b , $a \succ_i b$ means that the voter v_i prefers a to b . We denote the set of all complete orders over \mathcal{C} by $\mathcal{L}(\mathcal{C})$. Hence, $\mathcal{L}(\mathcal{C})^n$ denotes the set of all n -voters' preference profiles $(\succ_1, \dots, \succ_n)$.

A map $r : \mathcal{L}(\mathcal{C})^n \rightarrow \mathcal{C}$ is called a *voting rule*. Given a vote profile $\succ \in \mathcal{L}(\mathcal{C})^n$, we call $r(\succ)$ the *winner*. Note that in this paper, each election has a unique winner, and we ignore the possibility of ties. A voting rule is called *homogeneous* if it selects the winner solely based on the fraction of times each complete order from $\mathcal{L}(\mathcal{C})$ appears as a vote in the election. All the commonly used voting rules including the ones that are studied in this paper are homogeneous.

Given an election E , we can construct a weighted graph G_E called *weighted majority graph* from E . The set of vertices in G_E is the set of candidates in E . For any two candidates x and y , the weight on the edge (x, y) is $D_E(x, y) = N_E(x, y) - N_E(y, x)$, where $N_E(x, y)$ (respectively $N_E(y, x)$) is the number of voters who prefer x to y (respectively y to x). A candidate x is called the *Condorcet winner* in an election E if $D_E(x, y) > 0$ for every other candidate $y \neq x$. A voting rule is called *Condorcet consistent* if it selects the Condorcet winner as the winner of the election whenever it exists.

Some examples of common voting rules² are:

Positional scoring rules: A collection of m -dimensional vectors $\vec{s}_m = (\alpha_1, \alpha_2, \dots, \alpha_m) \in \mathbb{R}^m$ with $\alpha_1 \geq \alpha_2 \geq \dots \geq \alpha_m$ and $\alpha_1 > \alpha_m$ for every $m \in \mathbb{N}$ naturally defines a voting rule – a candidate gets score α_i from a vote if it is placed at the i^{th} position, and the score of a candidate is the sum of the scores it receives from all the votes. The winner is the candidate with maximum score.

Without loss of generality, we assume that for any score vector $\vec{\alpha}$, there exists a j such that $\alpha_j = 1$ and $\alpha_k = 0$ for all $k > j$. The vector α that is 1 in the first k coordinates and 0 otherwise gives the k -approval voting rule. 1-approval is called the *plurality* voting rule, and $(m-1)$ -approval is called the *veto* voting rule. The score vector $(m-1, m-2, \dots, 1, 0)$ gives the *Borda* voting rule.

²In all these rules, the possibilities of ties exist. If they do happen, we assume that some arbitrary but fixed tie breaking rule is applied.

Approval: In approval voting, each voter approves a subset of candidates. The winner is the candidate which is approved by the maximum number of voters.

Maximin: The maximin score of a candidate x is $\min_{y \neq x} D(x, y)$. The winner is the candidate with maximum maximin score.

Copeland ^{α} : The Copeland ^{α} score of a candidate x is $|\{y \neq x : D_E(x, y) > 0\}| + \alpha|\{y \neq x : D_E(x, y) = 0\}|$, where $\alpha \in [0, 1]$. The winner is the candidate with the maximum Copeland score.

Bucklin: A candidate x 's Bucklin score is the minimum number l such that more than half of the voters rank x in their top l positions. The winner is the candidate with lowest Bucklin score.

Plurality with runoff: The top two candidates according to plurality score are selected first. The pairwise winner of these two candidates is selected as the winner of the election. This rule is often called the *runoff* voting rule.

Single Transferable Vote: In Single Transferable Vote (STV), a candidate with least plurality score is dropped out of the election and its votes are transferred to the next preferred candidate. If two or more candidates receive least plurality score, then tie breaking rule is used. The candidate that remains after $(m-1)$ rounds is the winner.

Among the above voting rules, only the maximin and the Copeland voting rules are Condorcet consistent.

Given an election, the margin of victory of this election is:

Definition 2. Given a voting profile \succ , the margin of victory (MOV) is the smallest number of votes k such that the winner can be changed by changing k many votes in \succ , while keeping other votes unchanged.

Xia [32] showed that for most common voting rules (including all those mentioned above), when each voter votes i.i.d. according to a distribution on the candidates, the margin of victory is with high probability, either $\Theta(\sqrt{n})$ or $\Theta(n)$.

2.2 Statistical Distance Measures

Given a finite set X , a distribution μ on X is defined as a function $\mu : X \rightarrow [0, 1]$, such that $\sum_{x \in X} \mu(x) = 1$. The finite set X is called the base set of the distribution μ . We use the following distance measures among distributions in our work.

Definition 3. The KL divergence [21] and the Jensen-Shannon divergence [22] between two distributions μ_1 and μ_2 on X are defined as follows.

$$D_{KL}(\mu_1 || \mu_2) = \sum_{x \in X} \mu_1(x) \log \frac{\mu_1(x)}{\mu_2(x)}$$

$$JS(\mu_1, \mu_2) = \frac{1}{2} \left(D_{KL} \left(\mu_1 || \frac{\mu_1 + \mu_2}{2} \right) + D_{KL} \left(\mu_2 || \frac{\mu_1 + \mu_2}{2} \right) \right)$$

The Jensen-Shannon divergence has subsequently been generalized to measure the mutual distance among more than two distributions as follows.

Definition 4. Given n distributions μ_1, \dots, μ_n over the same base set, the generalized Jensen-Shannon divergence³ among them is:

$$JS(\mu_1, \dots, \mu_n) = \frac{1}{n} \sum_{i=1}^n D_{KL} \left(\mu_i || \frac{1}{n} \sum_{j=1}^n \mu_j \right)$$

³The generalized Jensen-Shannon divergence is often formulated with weights on each of the n distributions. The definition here puts equal weight on each distribution and is sufficient for our purposes.

2.3 Chernoff Bound

We repeatedly use the following concentration inequality:

Theorem 1. Let X_1, \dots, X_ℓ be a sequence of ℓ independent random variables in $[0, 1]$ (not necessarily identical). Let $S = \sum_i X_i$ and let $\mu = \mathbb{E}[S]$. Then, for any $0 \leq \delta \leq 1$:

$$\Pr[|S - \mu| \geq \delta \ell] < 2 \exp(-2\ell\delta^2)$$

and

$$\Pr[|S - \mu| \geq \delta \mu] < 2 \exp(-\delta^2 \mu/3)$$

The first inequality is called an additive bound and the second multiplicative.

3. RESULTS ON LOWER BOUNDS

Our lower bounds for the sample complexity of (ε, δ) -winner determination are derived from information-theoretic lower bounds for distinguishing distributions.

We start from the following basic observation. Let X be a random variable taking value 1 with probability $\frac{1}{2} - \varepsilon$ and 0 with probability $\frac{1}{2} + \varepsilon$; Y be a random variable taking value 1 with probability $\frac{1}{2} + \varepsilon$ and 0 with probability $\frac{1}{2} - \varepsilon$. Then, it is well-known that every algorithm needs $\Omega(\frac{1}{\varepsilon^2} \log \frac{1}{\delta})$ many samples to distinguish between X and Y with probability of making an error being at most δ [4, 9]. Immediately, we have:

Theorem 2. The sample complexity of the (ε, δ) -winner determination problem for the plurality voting rule is $\Omega(\frac{1}{\varepsilon^2} \log \frac{1}{\delta})$.

PROOF. Consider an election with two candidates a and b . Consider two vote distributions X and Y . In X , exactly $\frac{1}{2} + \varepsilon$ fraction of voters prefer a to b and thus a is the plurality winner of the election. In Y , exactly $\frac{1}{2} - \varepsilon$ fraction of voters prefer b to a and thus b is the plurality winner of the election. Also, the margin of victory of both the elections corresponding to the vote distributions X and Y is εn , since each vote change can change the plurality score of any candidate by at most one. Any (ε, δ) -winner determination algorithm for plurality will give us a distinguisher between the distributions X and Y with probability of error at most δ and hence will need $\Omega(\frac{1}{\varepsilon^2} \log \frac{1}{\delta})$ samples. \square

Theorem 2 immediately gives us the following corollary.

Corollary 1. Every (ε, δ) -winner determination algorithm needs $\Omega(\frac{1}{\varepsilon^2} \log \frac{1}{\delta})$ many samples for any voting rule which reduces to the plurality rule for two candidates. In particular, the lower bound holds for approval, scoring rules, maximin, Copeland, Bucklin, plurality with runoff, and STV voting rules.

We derive stronger lower bounds in terms of m by explicitly viewing the (ε, δ) -winner determination problem as a *statistical classification* problem. In this problem, we are given a black box that contains a distribution μ which is guaranteed to be one of ℓ known distributions μ_1, \dots, μ_ℓ . A *classifier* is a randomized oracle which has to determine the identity of μ , where each oracle call produces a sample from μ . At the end of its execution, the classifier announces a guess for the identity of μ , which has to be correct with probability at least $1 - \delta$. Using information-theoretic methods, Bar-Yossef [3] showed the following:

Lemma 1. The worst case sample complexity q of a classifier C for μ_1, \dots, μ_ℓ which does not make error with probability more than δ satisfies following.

$$q \geq \Omega\left(\frac{\log \ell}{JS(\mu_1, \dots, \mu_\ell)} \cdot (1 - \delta)\right)$$

The connection with our problem is the following. A set V of n votes on a candidate set \mathcal{C} generates a probability distribution μ_V on $\mathcal{L}(\mathcal{C})$, where $\mu_V(\succ)$ is proportional to the number of voters who voted \succ . Querying a random vote from V is then equivalent to sampling from the distribution μ_V . The margin of victory is proportional to the minimum statistical distance between μ_V and μ_W , over all the voting profiles W having a different winner than the winner of V .

Now, suppose we have m voting profiles V_1, \dots, V_m having different winners such that each V_i has margin of victory at least εn . Any (ε, δ) -winner determination algorithm must also be a statistical classifier for $\mu_{V_1}, \dots, \mu_{V_m}$ in the above sense. It then remains to construct such voting profiles for various voting rules which we do in the proof of the following theorem:

Theorem 3. Every (ε, δ) -winner determination algorithm needs $\Omega\left(\frac{\log m}{\varepsilon^2} \cdot (1 - \delta)\right)$ many samples for approval, Borda, Bucklin, and any Condorcet consistent voting rules, and $\Omega\left(\frac{\log(k+1)}{\varepsilon^2} \cdot (1 - \delta)\right)$ many samples for the k -approval voting rule.

PROOF. For each voting rules mentioned in the theorem, we will show m many distributions μ_1, \dots, μ_m on the votes with the following properties. Let V_i be an election where each vote v occurs exactly $\mu_i(v) \cdot n$ many times.

1. For every $i \neq j$, the winner in V_i is different from the winner in V_j .
2. For every i , the margin of victory of V_i is at least $\nu \varepsilon n$, for some constant $\nu \neq 0$.

The distributions for different voting rules are as follows. Let the candidate set be $\mathcal{C} = \{c_1, \dots, c_m\}$. Let $\mu = \frac{1}{m} \sum_{i=1}^m \mu_i$. For each case, it is easy to show that $D_{KL}(\mu_i || \mu) = O(\varepsilon^2)$. Hence, we only show that the margin of victory is as desired and thus the result follows from Lemma 1.

Approval, k -approval. Fix any arbitrary $M := k + 1$ many candidates c_1, \dots, c_M . For $i \in [M]$, we define a distribution μ_i on all k sized subsets of \mathcal{C} (for the k -approval voting rule, each vote is a k -sized subset of \mathcal{C}) as follows. Each k sized subset corresponds to top k candidates in a vote.

$$\mu_i(x) = \begin{cases} \frac{\varepsilon}{\binom{M-1}{k-1}} + \frac{1-\varepsilon}{\binom{M}{k}} & \text{if } c_i \in x \text{ and } x \subseteq \{c_1, \dots, c_M\} \\ \frac{1-\varepsilon}{\binom{M}{k}} & c_i \notin x \text{ and } x \subseteq \{c_1, \dots, c_M\} \\ 0 & \text{else} \end{cases}$$

The score of c_i in V_i is $n \left(\varepsilon + (1 - \varepsilon) \frac{\binom{M-1}{k-1}}{\binom{M}{k}} \right)$, the score of any

other candidate $c_j \in \{c_1, \dots, c_M\} \setminus \{c_i\}$ is $n(1 - \varepsilon) \frac{\binom{M-1}{k-1}}{\binom{M}{k}}$, and the score of the rest of the candidates is zero. Hence, the margin of victory is at least $\frac{\varepsilon n}{4}$, since each vote change can reduce the score of c_i by at most one and increase the score of any other candidate by at most one. This proves the result for the k -approval voting rule. We get the result for the approval voting rule from the fact that every k -approval election is also an approval election.

Borda, any Condorcet consistent voting rule. The score vector for the Borda voting rule which we use in this proof is $(m, m - 1, \dots, 1)$. For $i \in [m]$, we define a distribution μ_i on all possible linear orders over \mathcal{C} as follows.

$$\mu_i(x) = \begin{cases} \frac{\varepsilon}{(m-1)!} + \frac{1-\varepsilon}{m!} & \text{if } c_i \text{ is within top } \frac{m}{2} \text{ positions in } x. \\ \frac{1-\varepsilon}{m!} & \text{else} \end{cases}$$

The score of c_i in V_i is $\frac{mn}{2}(1 + \frac{\varepsilon}{2})$ whereas the score of any other candidate $c_j \neq c_i$ is $\frac{mn}{2}$. Hence, the margin of victory is at least $\frac{\varepsilon n}{8}$, since each vote change can reduce the score of c_i by at most m and increase the score of any other candidate by at most m . Also, in the weighted majority graph for the election V_i , $w(c_i, c_j) = \frac{\varepsilon n}{2}$. Hence, the margin of victory is at least $\frac{\varepsilon n}{4}$, since each vote change can change the weight of any edge in the weighted majority graph by at most two.

Bucklin. For $i \in [m]$, we define a distribution μ_i on all $\frac{m}{4}$ sized subsets of \mathcal{C} as follows. Each $\frac{m}{4}$ sized subset corresponds to the top $\frac{m}{4}$ candidates in a vote.

$$\mu_i(x) = \begin{cases} \frac{1-\varepsilon}{(\frac{m}{4}-1)} + \frac{\varepsilon}{(\frac{m}{4})} & \text{if } c_i \in x \\ \frac{\varepsilon}{(\frac{m}{4})} & \text{else} \end{cases}$$

The candidate c_i occurs within the top $\frac{m}{4}$ positions at least $n(1 - \frac{3\varepsilon}{4})$ many times, and any candidate $c_j \neq c_i$ occurs within the top $\frac{m}{4}$ positions at most $\frac{n}{3} - \frac{\varepsilon n}{12}$ many times. Hence, the margin of victory is at least $\frac{\varepsilon n}{6}$, since each vote change can change the number of time any particular candidate occurs within top $\frac{m}{4}$ positions by at most one. \square

4. RESULTS ON UPPER BOUNDS

In this section, we present the upper bounds on the sample complexity of the (ε, δ) -winner determination problem for various voting rules. The general framework for proving the upper bounds is as follows. For each voting rule, we first prove a useful structural property about the election when the margin of victory is known to be at least εn . Then, we sample a few votes uniformly at random to estimate either the score of the candidates for score based voting rules or weights of the edges in the weighted majority graph for other voting rules. Finally, appealing to the structural property that has been established, we argue that, the winner of the election on the sampled votes will be the same as the winner of the election, if we are able to estimate either the scores of the candidates or the weights of the edges in the weighted majority graph to a certain level of accuracy.

Before getting into specific voting rules, we prove a straightforward bound on the sample complexity for the (ε, δ) -winner determination problem for any voting rule.

Theorem 4. *There is a (ε, δ) -winner determination algorithm for every homogeneous voting rules with sample complexity $O(\frac{m!^2 \log \frac{m!}{\delta}}{\varepsilon^2})$.*

PROOF. We sample ℓ votes uniformly at random from the set of votes with replacement. Let X_i be an indicator random variable that is 1 exactly when x is the i 'th sample, and let $g(x)$ be the total number of voters whose vote is x . Define $\hat{g}(x) = \frac{n}{\ell} \sum_{i=1}^{\ell} X_i$. Using the Chernoff bound (Theorem 1), we have the following:

$$\Pr \left[|\hat{g}(x) - g(x)| > \frac{\varepsilon n}{2m!} \right] \leq 2 \cdot \exp \left(-\frac{\varepsilon^2 \ell}{2m!^2} \right)$$

By using the union bound, we have the following,

$$\Pr \left[\exists x \in \mathcal{L}(\mathcal{C}), |\hat{g}(x) - g(x)| > \frac{\varepsilon n}{2m!} \right] \leq 2m! \cdot \exp \left(-\frac{\varepsilon^2 \ell}{2m!^2} \right)$$

Since the margin of victory is εn and the voting rule is anonymous, the winner of the ℓ sample votes will be same as the winner of the election if $|\hat{g}(x) - g(x)| \leq \frac{\varepsilon n}{2m!}$ for every linear order $x \in \mathcal{L}(\mathcal{C})$. Hence, it is enough to take $\ell = O(m!^2 / \varepsilon^2 \cdot \log(m! / \delta))$. \square

4.1 Approval Voting Rule

We derive the upper bound on the sample complexity for the (ε, δ) -winner determination problem for the approval voting rule.

Lemma 2. *If $\text{MOV} \geq \varepsilon n$ and w be the winner of a approval election, then, $s(w) - s(x) \geq \varepsilon n$, for every candidate $x \neq w$, where $s(y)$ is the number of approvals that a candidate y receives.*

PROOF. Suppose there is a candidate $x \neq w$ such that $s(w) - s(x) < \varepsilon n$. Then there must exist $\varepsilon n - 1$ votes which does not approve the candidate x . We modify these votes to make it approve x . This makes w not the unique winner in the modified election. This contradicts the fact that the MOV is at least εn . \square

Theorem 5. *There is a (ε, δ) -winner determination algorithm for the approval voting rule with sample complexity $O(\frac{\log(m/\delta)}{\varepsilon^2})$.*

PROOF. Suppose w is the winner. We sample ℓ votes uniformly at random from the set of votes with replacement. For a candidate x , let X_i be a random variable indicating whether the i 'th vote sampled approved x . Define $\hat{s}(x) = \frac{n}{\ell} \sum_{i=1}^{\ell} X_i$. Then, by an argument analogous to the proof of Theorem 4, $\Pr[|\hat{s}(x) - s(x)| > \varepsilon n/2] \leq 2m \cdot \exp(-\varepsilon^2 \ell/2)$. Thus since $\text{MOV} \geq \varepsilon n$ and by Lemma 2, if we take $\ell = O(\frac{\log m/\delta}{\varepsilon^2})$, $\hat{s}(w)$ is greater than $\hat{s}(x)$ for all $x \neq w$. \square

4.2 Scoring Rules

Now, we move on to the scoring rules. Again, we first establish a structural consequence of having large MOV.

Lemma 3. *Suppose $\alpha = (\alpha_1, \dots, \alpha_m)$ be a normalized score vector and w is the winner of an election using scoring rule α with $\text{MOV} \geq \varepsilon n$. Then, $s(w) - s(x) \geq \alpha_1 \varepsilon n/2$ for every candidate $x \neq w$, where $s(w)$ and $s(x)$ denote the score of the candidates w and x respectively.*

PROOF. There must be at least εn many votes where w is preferred over x , since we can make x win the election by exchanging the positions of x and w in all these votes and $\text{MOV} \geq \varepsilon n$. Let v be a vote where w is preferred to x . Suppose we replace the vote v by another vote $v' = x \succ \text{others} \succ w$. We claim that this replacement reduces the current value of $s(w) - s(x)$ by at least α_1 . If we change $\varepsilon n/2$ such votes, then $s(w) - s(x)$ decreases by at least $\alpha_1 \varepsilon n/2$ but, at the same time, w must still be the winner after the vote changes because of the MOV condition. So, $s(w) - s(x) \geq \alpha_1 \varepsilon n/2$.

To prove the claim, suppose w and x were receiving a score of α_i and α_j respectively from the vote v . By replacing the vote v by v' , the current value of $s(w) - s(x)$ reduces by $\alpha_1 - \alpha_j + \alpha_i$, since $\alpha_m = 0$. Now, $\alpha_1 - \alpha_j + \alpha_i \geq \alpha_1$ since in the vote v , the candidate w is preferred over x and hence, $\alpha_j < \alpha_i$. This proves the result. \square

Theorem 6. *Suppose $\alpha = (\alpha_1, \dots, \alpha_m)$ be a normalized score vector. There is a (ε, δ) -winner determination algorithm for the α -scoring rule with sample complexity $O(\frac{\log(m/\delta)}{\varepsilon^2})$.*

PROOF. It is enough to show the result for the $(2\varepsilon, \delta)$ -winner determination problem. We sample ℓ votes uniformly at random from the set of votes with replacement. For a candidate x , define $X_i = \frac{\alpha_i}{\alpha_1}$ if x gets a score of α_i from the i th sample vote, and let $\hat{s}(x) = \frac{n\alpha_1}{\ell} \sum_{i=1}^{\ell} X_i$. Now, using Chernoff bound (Theorem 1), we have:

$$\Pr[|\hat{s}(x) - s(x)| \geq \alpha_1 \varepsilon n/4] \leq 2 \exp \left(-\frac{\varepsilon^2 \ell}{2} \right)$$

The rest of the proof follows from an argument analogous to the proof of Theorem 5 using Lemma 3. \square

From Theorem 6, we have a (ε, δ) -winner determination algorithm for the k -approval voting rule which needs $O(\frac{\log(m/\delta)}{\varepsilon^2})$ many samples for any k . This is tight by Theorem 3 when $k = cm$ for some $c \in (0, 1)$.

When $k = o(m)$, we have a lower bound of $\Omega(\frac{1}{\varepsilon^2} \log \frac{1}{\delta})$ for the k -approval voting rule (see Corollary 1). We show next that this lower bound is also tight for the k -approval voting rule when $k = o(m)$. Before embarking on the proof of the above fact, we prove the following lemma which will be crucially used.

Lemma 4. *Let $f : \mathbb{R} \rightarrow \mathbb{R}$ be a function defined by $f(x) = e^{-\frac{\lambda}{x}}$. Then,*

$$f(x) + f(y) \leq f(x + y), \text{ for } x, y > 0, \frac{\lambda}{x + y} > 2, x < y$$

PROOF. For the function $f(x)$, we have following.

$$\begin{aligned} f(x) &= e^{-\frac{\lambda}{x}} \\ \Rightarrow f'(x) &= \frac{\lambda}{x^2} e^{-\frac{\lambda}{x}} \\ \Rightarrow f''(x) &= \frac{\lambda^2}{x^4} e^{-\frac{\lambda}{x}} - \frac{2\lambda}{x^3} e^{-\frac{\lambda}{x}} \end{aligned}$$

Hence, for $x, y > 0, \frac{\lambda}{x+y} > 2, x < y$ we have $f''(x), f''(y), f''(x+y) > 0$. This implies following for $x < y$ and an infinitesimal positive δ .

$$\begin{aligned} f'(x) &\leq f'(y) \\ \Rightarrow \frac{f(x - \delta) - f(x)}{\delta} &\geq \frac{f(y) - f(y - \delta)}{\delta} \\ \Rightarrow f(x) + f(y) &\leq f(x - \delta) + f(y + \delta) \\ \Rightarrow f(x) + f(y) &\leq f(x + y) \end{aligned}$$

□

Theorem 7. *There is a (ε, δ) -winner determination algorithm for the k -approval voting rule with sample complexity $O(\frac{\log(\frac{k}{\delta})}{\varepsilon^2})$.*

PROOF. It is enough to show the result for the $(2\varepsilon, \delta)$ -winner determination problem. We sample ℓ votes uniformly at random from the set of votes with replacement. For a candidate x , let X_i be a random variable indicating whether x is among the top k candidates for the i^{th} vote sample. Define $\hat{s}(x) = \frac{n}{\ell} \sum_{i=1}^{\ell} X_i$, and let $s(x)$ be the actual score of x . Then by the multiplicative Chernoff bound (Theorem 1), we have:

$$\Pr[|\hat{s}(x) - s(x)| > \varepsilon n] \leq 2 \exp\left(-\frac{\varepsilon^2 \ell n}{3s(x)}\right)$$

By union bound, we have the following,

$$\begin{aligned} &\Pr[\exists x \in \mathcal{C}, |\hat{s}(x) - s(x)| > \varepsilon n] \\ &\leq \sum_{x \in \mathcal{C}} 2 \exp\left(-\frac{\varepsilon^2 \ell n}{3s(x)}\right) \\ &\leq 2k \exp(-\varepsilon^2 \ell / 3) \end{aligned}$$

Let the candidate w be the winner of the election. The second inequality in the above derivation follows from the fact that, the function $\sum_{x \in \mathcal{C}} \exp\left(-\frac{\varepsilon^2 \ell n}{3s(x)}\right)$ is maximized in the domain, defined by the constraint: for every candidate $x \in \mathcal{C}$, $s(x) \in [0, n]$ and $\sum_{x \in \mathcal{C}} s(x) = kn$, by setting $s(x) = n$ for every $x \in \mathcal{C}'$ and $s(y) = 0$ for every $y \in \mathcal{C} \setminus \mathcal{C}'$, for any arbitrary subset $\mathcal{C}' \subset \mathcal{C}$ of cardinality k (due to Lemma 4). The rest of the proof follows by an argument analogous to the proof of Theorem 4 using Lemma 3. □

Notice that, the sample complexity upper bound in Theorem 7 is independent of m for the plurality voting rule. Theorem 7 in turn implies the following Corollary which we consider to be of independent interest.

Corollary 2. *There is an algorithm to estimate the ℓ_∞ norm $\ell_\infty(\mu)$ of a distribution μ within an additive factor of ε by querying only $O(\frac{1}{\varepsilon^2} \log \frac{1}{\delta})$ many samples, if we are allowed to get i.i.d. samples from the distribution μ .*

Such a statement seems to be folklore in the statistics community [15]. Recently in an independent and nearly simultaneous work, Waggoner [31] obtained a sharp bound of $\frac{4}{\varepsilon^2} \log(\frac{1}{\delta})$ for the sample complexity in Corollary 2.

4.3 Maximin Voting Rule

We now turn our attention to the maximin voting rule. The idea is to sample enough number of votes such that we are able to estimate the weights of the edges in the weighted majority graph with certain level of accuracy which in turn leads us to predict winner.

Lemma 5. *Suppose $\text{MOV} \geq \varepsilon n$ and w be the winner of a maximin election. Then, $s(w) - s(x) \geq \varepsilon n$, for every candidate $x \neq w$, where $s(\cdot)$ is the maximin score.*

PROOF. Let w be the winner and x be any arbitrary candidate other than w . Suppose, for contradiction, $s(w) - s(x) < \varepsilon n$. Suppose y be a candidate such that $N(w, y) = s(w)$. Now there exist at least $\varepsilon n - 1$ votes as below.

$$c_1 \succ \dots \succ w \succ \dots \succ y \succ \dots \succ c_{m-2}$$

We replace $\varepsilon n - 1$ of such votes by the votes as below.

$$c_1 \succ \dots \succ y \succ \dots \succ c_{m-2} \succ w$$

This makes the maximin score of w less than the maximin score of x . This contradicts the assumption that $\text{MOV} \geq \varepsilon n$. □

Theorem 8. *There is a (ε, δ) -winner determination algorithm for the maximin voting rule with sample complexity $O(\frac{\log(\frac{m}{\delta})}{\varepsilon^2})$.*

PROOF. Let x and y be any two arbitrary candidates. We sample ℓ votes uniformly at random from the set of votes with replacement. Let X_i be a random variable defined as follows.

$$X_i = \begin{cases} 1, & \text{if } x \succ y \text{ in the } i^{\text{th}} \text{ sample} \\ -1, & \text{else} \end{cases}$$

Define $\hat{D}(x, y) = \frac{n}{\ell} \sum_{i=1}^{\ell} X_i$. We estimate $\hat{D}(x, y)$ within the closed ball of radius $\varepsilon n/2$ around $D(x, y)$ for every candidates $x, y \in \mathcal{C}$ and the rest of the proof follows from by an argument analogous to the proof of Theorem 5 using Lemma 5. □

4.4 Copeland $^\alpha$ Voting Rule

Now, we move on to the Copeland $^\alpha$ voting rule. The approach is similar to the maximin voting rule. However, it turns out that we need to estimate the edge weights of the weighted majority graph more accurately for the Copeland $^\alpha$ voting rule. Xia introduced the brilliant quantity called the *relative margin of victory* (see Section 5.1 in [32]) which will be used crucially for showing sample complexity upper bound for the Copeland $^\alpha$ voting rule. Given an election, a candidate $x \in \mathcal{C}$, and an integer (may be negative also) t , $s'_t(V, x)$ is defined as follows.

$$\begin{aligned} s'_t(V, x) &= |\{y \in \mathcal{C} : y \neq x, D_\varepsilon(y, x) < 2t\}| \\ &\quad + \alpha |\{y \in \mathcal{C} : y \neq x, D_\varepsilon(y, x) = 2t\}| \end{aligned}$$

For every two distinct candidates x and y , the relative margin of victory, denoted by $RM(x, y)$, between x and y is defined as the minimum integer t such that, $s'_{-t}(V, x) \leq s'_t(V, y)$. Let w be the winner of the election \mathcal{E} . We define a quantity $\Gamma(\mathcal{E})$ to be $\min_{x \in C \setminus \{w\}} \{RM(w, x)\}$. Notice that, given an election \mathcal{E} , $\Gamma(\mathcal{E})$ can be computed in polynomial amount of time. Now we have the following lemma.

Lemma 6. Suppose $MOV \geq \varepsilon n$ and w be the winner of a Copeland $^\alpha$ election. Then, $RM(w, x) \geq \frac{\varepsilon n}{2(\lceil \log m \rceil + 1)}$, for every candidate $x \neq w$.

PROOF. Follows from Theorem 11 in [32]. \square

Theorem 9. There is a (ε, δ) -winner determination algorithm for Copeland $^\alpha$ voting rule with sample complexity $O(\frac{\log^3 \frac{m}{\delta}}{\varepsilon^2})$.

PROOF. Let x and y be any two arbitrary candidates and w the Copeland $^\alpha$ winner of the election. We estimate $D(x, y)$ within the closed ball of radius $\frac{\varepsilon n}{5(\lceil \log m \rceil + 1)}$ around $D(x, y)$ for every candidates $x, y \in C$ in a way analogous to the proof of Theorem 8. This needs $O(\frac{\log^3 \frac{m}{\delta}}{\varepsilon^2})$ many samples. The rest of the proof follows from Lemma 6 by an argument analogous to the proof of Theorem 4. \square

4.5 Bucklin Voting Rule

For the Bucklin voting rule, we will estimate how many times each candidate occurs within the first k position for every $k \in [m]$. This eventually leads us to predict the winner of the election due to the following lemma.

Lemma 7. Suppose MOV of a Bucklin election be at least εn . Let w be the winner of the election and x be any arbitrary candidate other than w . Suppose

$$b_w = \min_i \{i : w \text{ is within top } i \text{ places in at least } \frac{n}{2} + \frac{\varepsilon n}{3} \text{ votes}\}$$

$$b_x = \min_i \{i : x \text{ is within top } i \text{ places in at least } \frac{n}{2} - \frac{\varepsilon n}{3} \text{ votes}\}$$

Then, $b_w < b_x$.

PROOF. We prove it by contradiction. So, assume $b_w \geq b_x$. Now by changing $\frac{\varepsilon n}{3}$ votes, we can make the Bucklin score of w to be at least b_w . By changing another $\frac{\varepsilon n}{3}$ votes, we can make the Bucklin score of x to be at most b_x . Hence, by changing $\frac{2\varepsilon n}{3}$ votes, it is possible not to make w the unique winner which contradicts the fact that the MOV is at least εn . \square

Theorem 10. There is a (ε, δ) -winner determination algorithm for Bucklin voting rule with sample complexity $O(\frac{\log \frac{m}{\delta}}{\varepsilon^2})$.

PROOF. Let x be any arbitrary candidate and $1 \leq k \leq m$. We sample l votes uniformly at random from the set of votes with replacement. Let X_i be a random variable defined as follows.

$$X_i = \begin{cases} 1, & \text{if } x \text{ is within top } k \text{ places in } i^{th} \text{ sample} \\ 0, & \text{else} \end{cases}$$

Let $\hat{s}_k(x)$ be the estimate of the number of times the candidate x has been placed within top k positions. That is, $\hat{s}_k(x) = \frac{n}{l} \sum_{i=1}^l X_i$. Let $s_k(x)$ be the number of times the candidate x been placed in top k positions. Clearly, $E[\hat{s}_k(x)] = \frac{n}{l} \sum_{i=1}^l E[X_i] = s_k(x)$. We estimate $\hat{s}_k(x)$ within the closed ball of radius $\varepsilon n/2$ around $s_k(x)$ for every candidate $x \in C$ and every integer $k \in [m]$, and the rest of the proof follows from by an argument analogous to the proof of Theorem 5 using Lemma 7. \square

4.6 Plurality with Runoff Voting Rule

Now, we move on to the plurality with runoff voting rule. In this case, we first estimate the plurality score of each of the candidates. In the next round, we estimate the pairwise margin of victory of the two candidates that qualifies to the second round.

Lemma 8. Suppose $MOV \geq \varepsilon n$, and w and r be the winner and runner up of a plurality with runoff election respectively, and x be any arbitrary candidate other than w and r . Then, following holds. Let $s(\cdot)$ denote plurality score of candidates. Then following holds.

1. $D(w, r) > 2\varepsilon n$.
2. For every candidate $x \in C \setminus \{w, r\}$, $2s(w) > s(x) + s(r) + \varepsilon n$.
3. If $s(x) > s(r) - \frac{\varepsilon n}{2}$, then $D(w, x) > \frac{\varepsilon n}{2}$.

PROOF. If the first property does not hold, then by changing εn votes, we can make r winner. If the second property does not hold, then by changing εn votes, we can make both x and r qualify to the second round. If the third property does not hold, then by changing $\frac{\varepsilon n}{2}$ votes, the candidate x can be sent to the second round of the runoff election. By changing another $\frac{\varepsilon n}{2}$ votes, x can be made to win the election. This contradicts the MOV assumption. \square

Theorem 11. There is a (ε, δ) -winner determination algorithm for the plurality with runoff voting rule with sample complexity $O(\frac{\log \frac{1}{\delta}}{\varepsilon^2})$.

PROOF. Let x be any arbitrary candidate. We sample l votes uniformly at random from the set of votes with replacement. Let, X_i be a random variable defined as follows.

$$X_i = \begin{cases} 1, & \text{if } x \text{ is at first position in the } i^{th} \text{ sample} \\ 0, & \text{else} \end{cases}$$

The estimate of the plurality score of x be $\hat{s}(x)$. Then $\hat{s}(x) = \frac{n}{l} \sum_{i=1}^l X_i$. Let $s(x)$ be the actual plurality score of x . Then we have following,

$$E[X_i] = \frac{s(x)}{n}, E[\hat{s}(x)] = \frac{n}{l} \sum_{i=1}^l E[X_i] = s(x)$$

By Chernoff bound, we have the following,

$$\Pr[|\hat{s}(x) - s(x)| > \varepsilon n] \leq \frac{2}{\exp\{\varepsilon^2 l n / 3s(x)\}}$$

By union bound, we have the following,

$$\begin{aligned} \Pr[\exists x \in C, |\hat{s}(x) - s(x)| > \varepsilon n] &\leq \sum_{x \in C} \frac{2}{\exp\{\varepsilon^2 l n / 3s(x)\}} \\ &\leq \frac{2}{\exp\{\varepsilon^2 l / 3\}} \end{aligned}$$

The last line follows from Lemma 4. Notice that, we do not need the random variables $\hat{s}(x)$ and $\hat{s}(y)$ to be independent for any two candidates x and y . Hence, we can use the same l sample votes to estimate $\hat{s}(x)$ for every candidate x .

Now, let y and z be the two candidates that go to the second round.

$$Y_i = \begin{cases} 1, & \text{if } y \succ z \text{ in the } i^{th} \text{ sample} \\ -1, & \text{else} \end{cases}$$

The estimate of $D(y, z)$ be $\hat{D}(y, z)$. Then $\hat{D}(y, z) = \frac{n}{l} \sum_{i=1}^l Y_i$. Then we have following,

$$E[Y_i] = \frac{D(y, z)}{n}, E[\hat{D}(y, z)] = \frac{n}{l} \sum_{i=1}^l E[Y_i] = D(y, z)$$

By Chernoff bound, we have the following,

$$\Pr[|\hat{D}(y, z) - D(y, z)| > \varepsilon n] \leq \frac{2}{\exp\{\varepsilon^2 l/3\}}$$

Let A be the event that $\forall x \in \mathcal{C}$, $|\hat{s}(x) - s(x)| \leq \varepsilon n$ and $|\hat{D}(y, z) - D(y, z)| \leq \varepsilon n$. Now we have,

$$\Pr[A] \geq 1 - \left(\frac{2}{\exp\{\varepsilon^2 l/3\}} + \frac{2}{\exp\{\varepsilon^2 l/3\}} \right)$$

Since we do not need independence among the random variables $\hat{s}(a)$, $\hat{s}(b)$, $\hat{D}(w, x)$, $\hat{D}(y, z)$ for any candidates a, b, w, x, y , and z , we can use the same l sampled votes. Now, from Lemma 8, if $|\hat{s}(x) - s(x)| \leq \frac{\varepsilon n}{5}$ for every candidate x and $|\hat{D}(y, z) - D(y, z)| \leq \frac{\varepsilon n}{5}$ for every candidates y and z , then the plurality with runoff winner of the sampled votes coincides with the actual runoff winner. The above event happens with probability at least $1 - \delta$ by choosing an appropriate $l = O(\frac{\log \frac{1}{\delta}}{\varepsilon^2})$. \square

4.7 STV Voting Rule

Now we move on the STV voting rule. The following lemma provides an upper bound on the number of votes that need to be changed to make some arbitrary candidate win the election. More specifically, given a sequence of m candidates $\{x_i\}_{i=1}^m$ with x_m not being the winner, the lemma below proves an upper bound on the number of number of votes that need to be modified such that the candidate x_i gets eliminated at the i^{th} round in the STV voting rule.

Lemma 9. Suppose \mathcal{V} be a set of votes and w be the winner of a STV election. Consider the following chain with candidates $x_1 \neq x_2 \neq \dots \neq x_m$ and $x_m \neq w$.

$$\mathcal{C} \supset \mathcal{C} \setminus \{x_1\} \supset \mathcal{C} \setminus \{x_1, x_2\} \supset \dots \supset \{x_m\}$$

Let $s_{\mathcal{V}}(A, x)$ be the plurality score of a candidate x when all the votes in \mathcal{V} are restricted to the set of candidates $A \subset \mathcal{C}$. Let us define $\mathcal{C}_{-i} = \mathcal{C} \setminus \{x_1, \dots, x_i\}$ and $s_{\mathcal{V}}^*(A) := \min_{x \in A} \{s_{\mathcal{V}}(A, x)\}$. Then, we have the following.

$$\sum_{i=0}^{m-1} (s_{\mathcal{V}}(\mathcal{C}_{-i}, x_{i+1}) - s_{\mathcal{V}}^*(\mathcal{C}_{-i})) \geq MOV$$

PROOF. We will show that by changing $\sum_{i=0}^{m-1} (s_{\mathcal{V}}(\mathcal{C}_{-i}, x_{i+1}) - s_{\mathcal{V}}^*(\mathcal{C}_{-i}))$ votes, we can make the candidate x_m winner. If x_1 minimizes $s_{\mathcal{V}}(\mathcal{C}, x)$ over $x \in \mathcal{C}$, then we do not change anything and define $\mathcal{V}_1 = \mathcal{V}$. Otherwise, there exist $s_{\mathcal{V}}(\mathcal{C}, x_1) - s_{\mathcal{V}}^*(\mathcal{C})$ many votes of following type.

$$x_1 \succ a_1 \succ a_2 \succ \dots \succ a_{m-1}, a_i \in \mathcal{C}, \forall 1 \leq i \leq m-1$$

We replace $s_{\mathcal{V}}(\mathcal{C}, x_1) - s_{\mathcal{V}}^*(\mathcal{C})$ many votes of the above type by the votes as follows.

$$a_1 \succ x_1 \succ a_2 \succ \dots \succ a_{m-1}$$

Let us call the new set of votes by \mathcal{V}_1 . We claim that, $s_{\mathcal{V}}(\mathcal{C} \setminus \{x_1\}, x) = s_{\mathcal{V}_1}(\mathcal{C} \setminus \{x_1\}, x)$ for every candidate $x \in \mathcal{C} \setminus \{x_1\}$. Fix any arbitrary candidate $x \in \mathcal{C} \setminus \{x_1\}$. The votes in \mathcal{V}_1 that are same as in \mathcal{V} contributes same quantity to both side of the equality. Let v be a vote that has been changed as described above. If $x = a_1$ then, the vote v contributes one to both sides of the equality. If $x \neq a_1$, then the vote contributes zero to both sides of the equality. Hence, we have the claim. We repeat this process for $(m-1)$ times. Let \mathcal{V}_i be the set of votes after the candidate x_i gets eliminated. Now, in the above argument, by replacing \mathcal{V} by \mathcal{V}_{i-1} , \mathcal{V}_1 by \mathcal{V}_i , the candidate

set \mathcal{C} by $\mathcal{C} \setminus \{x_1, \dots, x_{i-1}\}$, and the candidate x_1 by the candidate x_i , we have the following.

$$s_{\mathcal{V}_{i-1}}(\mathcal{C}_{-i}, x) = s_{\mathcal{V}_i}(\mathcal{C}_{-i}, x) \forall x \in \mathcal{C} \setminus \{x_1, \dots, x_i\}$$

Hence, we have the following.

$$s_{\mathcal{V}}(\mathcal{C}_{-i}, x) = s_{\mathcal{V}_i}(\mathcal{C}_{-i}, x) \forall x \in \mathcal{C} \setminus \{x_1, \dots, x_i\}$$

In the above process, the total number of votes that are changed is $\sum_{i=0}^{m-1} (s_{\mathcal{V}}(\mathcal{C}_{-i}, x_{i+1}) - s_{\mathcal{V}}^*(\mathcal{C}_{-i}))$. \square

Theorem 12. There is a (ε, δ) -winner determination algorithm for the STV voting rule with sample complexity $O(\frac{m^2(m+\log \frac{1}{\delta})}{\varepsilon^2})$.

PROOF. We sample l votes uniformly at random from the set of votes with replacement and output the STV winner of those l votes say w' as the winner of the election. Let, w be the winner of the election. We will show that there exist $l = O(\frac{m^2(m+\log \frac{1}{\delta})}{\varepsilon^2})$ for which $w = w'$ with probability at least $1 - \delta$. Let A be an arbitrary subset of candidates and x be any candidate in A . Let us define a random variables X_i , $1 \leq i \leq l$ as follows.

$$X_i = \begin{cases} 1, & \text{if } x \text{ is at top } i^{th} \text{ sample when restricted to } A \\ 0, & \text{else} \end{cases}$$

Define another random variable $\hat{s}_{\mathcal{V}}(A, x) := \sum_{i=1}^l X_i$. Then we have, $E[\hat{s}_{\mathcal{V}}(A, x)] = s_{\mathcal{V}}(A, x)$. Now, using Chernoff bound, we have the following,

$$\Pr[|\hat{s}_{\mathcal{V}}(A, x) - s_{\mathcal{V}}(A, x)| > \frac{\varepsilon n}{m}] \leq \frac{2}{\exp\{\frac{\varepsilon^2 l}{3m^2}\}}$$

Let E be the event that $\exists A \subset \mathcal{C}$ and $\exists x \in A$, $|\hat{s}_{\mathcal{V}}(A, x) - s_{\mathcal{V}}(A, x)| > \frac{\varepsilon n}{m}$. By union bound, we have,

$$\Pr[\bar{E}] \geq 1 - \frac{m2^{m+1}}{\exp\{\frac{\varepsilon^2 l}{3m^2}\}}$$

The rest of the proof follows by an argument analogous to the proof of Theorem 4 using Lemma 9. \square

5. CONCLUSION

In this work, we introduced the (ε, δ) -winner determination problem and showed (often tight) bounds for the sample complexity for many common voting rules. Besides closing the remaining gaps in the bounds, here are a few open directions to pursue in the future:

- Is there an axiomatic characterization of the voting rules for which the sample complexity is independent of m and n ? We note that a similar problem in graph property testing was the subject of intense study [1, 7].
- Specifically for scoring rules, is the sample complexity determined by some natural property of the score vector, such as its sparsity?
- Is it worthwhile for the algorithm to elicit only part of the vote from each sampled voter instead of the full vote? As mentioned in the Introduction, vote elicitation is a well-trodden area, but as far as we know, it has not been studied how assuming a margin of victory can change the number of queries.
- How can knowledge of a social network on the voters be used to minimize the number of samples made? Some initial progress in this direction has been made by Dhamal and Narahari [12] and by Agrawal and Devanur (private communication).

REFERENCES

- [1] N. Alon, E. Fischer, I. Newman, and A. Shapira. A combinatorial characterization of the testable graph properties: it's all about regularity. In *Proceedings of the 38th Annual ACM Symposium on Theory of Computing (STOC)*, pages 251–260, 2006.
- [2] Y. Bachrach, N. Betzler, and P. Faliszewski. Probabilistic possible winner determination. In *International Conference on Artificial Intelligence (AAAI)*, volume 10, pages 697–702, 2010.
- [3] Z. Bar-Yossef. Sampling lower bounds via information theory. In *Proceedings of the 35th Annual ACM Symposium on Theory of Computing (STOC)*, pages 335–344. ACM, 2003.
- [4] Z. Bar-Yossef, R. Kumar, and D. Sivakumar. Sampling algorithms: lower bounds and applications. In *Proceedings of the 33rd Annual ACM Symposium on Theory of Computing (STOC)*, pages 266–275. ACM, 2001.
- [5] J. Bartholdi III, C. A. Tovey, and M. A. Trick. Voting schemes for which it can be difficult to tell who won the election. *Social Choice and Welfare (SCW)*, 6(2):157–165, 1989.
- [6] P. Boldi, F. Bonchi, C. Castillo, and S. Vigna. Voting in social networks. In *Proceedings of the 18th ACM Conference on Information and Knowledge Management*, pages 777–786. ACM, 2009.
- [7] C. Borgs, J. T. Chayes, L. Lovász, V. T. Sós, B. Szegedy, and K. Vesztegombi. Graph limits and parameter testing. In *Proceedings of the 38th Annual ACM Symposium on Theory of Computing (STOC)*, pages 261–270, 2006.
- [8] C. Boutilier, J. Lang, J. Oren, and H. Palacios. Robust winners and winner determination policies under candidate uncertainty. In *Proceedings of the International Conference on Artificial Intelligence (AAAI)*, 2014.
- [9] R. Canetti, G. Even, and O. Goldreich. Lower bounds for sampling algorithms for estimating the average. *Information Processing Letters (IPL)*, 53(1):17–25, 1995.
- [10] V. Conitzer. Eliciting single-peaked preferences using comparison queries. *Journal of Artificial Intelligence Research (JAIR)*, 35:161–191, 2009.
- [11] V. Conitzer and T. Sandholm. Vote elicitation: Complexity and strategy-proofness. In *Eighteenth International Conference on Artificial Intelligence (AAAI)*, pages 392–397, Menlo Park, CA, USA, 2002. American Association for Artificial Intelligence.
- [12] S. Dhamal and Y. Narahari. Scalable preference aggregation in social networks. In *First AAAI Conference on Human Computation and Crowdsourcing (HCOMP)*, 2013.
- [13] N. Ding and F. Lin. Voting with partial information: Minimal sets of questions to decide an outcome. In *Proc. Fourth International Workshop on Computational Social Choice (COMSOC-2012)*, Kraków, Poland, 2012.
- [14] J. A. Doucette, K. Larson, and R. Cohen. Approximate winner selection in social choice with partial preferences. In *Proceedings of the 12th International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*. International Foundation for Autonomous Agents and Multiagent Systems, 2014.
- [15] A. Dvoretzky, J. Kiefer, and J. Wolfowitz. Asymptotic minimax character of the sample distribution function and of the classical multinomial estimator. *The Annals of Mathematical Statistics*, pages 642–669, 1956.
- [16] E. Ephrati and J. Rosenschein. The Clarke tax as a consensus mechanism among automated agents. In *Proceedings of the Ninth International Conference on Artificial Intelligence (AAAI)*, pages 173–178, 1991.
- [17] D. Falik, R. Meir, and M. Tennenholtz. On coalitions and stable winners in plurality. In *8th International Workshop on Internet and Network Economics (WINE)*, pages 256–269. Springer, 2012.
- [18] N. Hazon, Y. Aumann, S. Kraus, and M. Wooldridge. Evaluation of election outcomes under uncertainty. In *Proceedings of the 7th International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, pages 959–966. International Foundation for Autonomous Agents and Multiagent Systems, 2008.
- [19] E. Hemaspaandra, L. A. Hemaspaandra, and J. Rothe. Exact analysis of dodgson elections: Lewis carroll's 1876 voting system is complete for parallel access to np. *Journal of the ACM (JACM)*, 44(6):806–825, 1997.
- [20] E. Hemaspaandra, H. Spakowski, and J. Vogel. The complexity of kemeny elections. *Theoretical Computer Science (TCS)*, 349(3):382–391, 2005.
- [21] S. Kullback and R. A. Leibler. On information and sufficiency. *The Annals of Mathematical Statistics*, pages 79–86, 1951.
- [22] J. Lin. Divergence measures based on the shannon entropy. *IEEE Transactions on Information Theory*, 37(1):145–151, 1991.
- [23] T. Lu and C. Boutilier. Robust approximation and incremental elicitation in voting protocols. In *Proceedings of the 22nd International Joint Conference on Artificial Intelligence (IJCAI)*, volume 22, page 287, 2011.
- [24] T. Lu and C. Boutilier. Vote elicitation with probabilistic preference models: Empirical estimation and cost tradeoffs. In *Algorithmic Decision Theory*, pages 135–149. Springer, 2011.
- [25] T. Lu and C. Boutilier. Multi-winner social choice with incomplete preferences. In *Proceedings of the 23rd International Joint Conference on Artificial Intelligence (IJCAI)*, pages 263–270. AAAI Press, 2013.
- [26] J. Oren, Y. Filmus, and C. Boutilier. Efficient vote elicitation under candidate uncertainty. In *Proceedings of the 23rd International Joint Conference on Artificial Intelligence (IJCAI)*, pages 309–316. AAAI Press, 2013.
- [27] D. M. Pennock, E. Horvitz, and C. L. Giles. Social choice theory and recommender systems: Analysis of the axiomatic foundations of collaborative filtering. In *Proceedings of the 17th International Conference on Artificial Intelligence (AAAI)*, 2000.
- [28] M. A. Rodriguez, D. J. Steinbock, J. H. Watkins, C. Gershenson, J. Bollen, V. Grey, and B. Degraf. Smartocracy: Social networks for collective decision making. In *40th Annual Hawaii International Conference on Systems Science (HICSS'07)*. Waikoloa, pages 90–90. IEEE, 2007.
- [29] D. Ron. Property testing. *Combinatorial Optimization-Dordrecht*, 9(2):597–643, 2001.
- [30] D. Shiryaev, L. Yu, and E. Elkind. On elections with robust winners. In *Proceedings of the 12th International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*,

- pages 415–422. International Foundation for Autonomous Agents and Multiagent Systems, 2013.
- [31] B. Waggoner. Lp testing and learning of discrete distributions. In *Proceedings of the 2015 Conference on Innovations in Theoretical Computer Science, ITCS '15*, pages 347–356, New York, NY, USA, 2015. ACM.
- [32] L. Xia. Computing the margin of victory for various voting rules. In *Proceedings of the 13th ACM Conference on Electronic Commerce (EC)*, pages 982–999. ACM, 2012.